

NO SANE COMPILER WOULD OPTIMIZE ATOMICS

JF Bastien <[@jfbastien](https://twitter.com/jfbastien)>

Compiler engineer on Chrome

false

DISCLAIMER



A black and white photograph showing a close-up of a person's hands working on a loom. The hands are positioned over a dark, textured surface, likely the loom frame. In the foreground, there is a large, soft pile of various colored yarns, including shades of red, orange, yellow, green, blue, and purple, some of which appear to be spilling onto the surface below. The background is out of focus, showing more of the loom and possibly other craft materials.

WHY DO I CARE?

A dark, moody photograph of a person's face in profile, looking down and to the side. Their hand is resting near their chin, partially obscuring their mouth. The lighting is dramatic, with strong highlights and shadows.

QUICK REVIEW

C++11 added

- A memory model
- Threads
- Classes to communicate between threads

Threads didn't exist before C++11

SEQUENTIAL CONSISTENCY

Naïve model: memory accesses are simply interleaved

Hardware doesn't work that way

Overly restrictive

SC-DRF

Data race:

- Two accesses to the same *memory location* by different threads are *not ordered*
- At least one of them stores to the memory location
- At least one of them is not a synchronization action

Data races are UB

MEMORY LOCATIONS

Either an object of *scalar type* or a maximal sequence of adjacent bit-fields all having non-zero width.

Arithmetic types, enumeration types, pointer types, pointer to member types, `std::nullptr_t`, and cv-qualified versions of these types are collectively called *scalar types*.

ORDERING MEMORY LOCATIONS

Two memory operations are ordered if they cannot occur simultaneously

ORDERING MEMORY LOCATIONS

An evaluation A *happens before* an evaluation B if:

- A is *sequenced before* B, or
- A *inter-thread happens before* B

SEQUENCED BEFORE

An asymmetric, transitive, pair-wise relation between evaluations executed by a single thread, which induces a partial order among those evaluations. Given any two evaluations A and B, if A is *sequenced before* B, then the execution of A shall precede the execution of B. If A is not *sequenced before* B and B is not *sequenced before* A, then A and B are *unsequenced*.

INTER-THREAD HAPPENS BEFORE

- A *synchronizes with* B, or
- A is *dependency-ordered before* B, or
- for some evaluation X
 - A *synchronizes with* X and X is *sequenced before* B, or
 - A is *sequenced before* X and X *inter-thread happens before* B, or
 - A *inter-thread happens before* X and X *inter-thread happens before* B.

SYNCHRONIZES WITH

An atomic operation A that performs a release operation on an atomic object M *synchronizes with* an atomic operation B that performs an acquire operation on M and takes its value from any side effect in the release sequence headed by A.

DEPENDENCY-ORDERED BEFORE

- A performs a release operation on an atomic object M, and, in another thread, B performs a consume operation on M and reads a value written by any side effect in the release sequence headed by A, or
- for some evaluation X, A is *dependency-ordered before* X and X carries a dependency to B.

[Note: The relation “is dependency-ordered before” is analogous to “synchronizes with”, but uses release/consume in place of release/acquire. – end note]

SYNCHRONIZATION OPERATION

- On one or more memory locations: consume operation, an acquire operation, a release operation, or both an acquire and release operation.
- Without an associated memory location: acquire fence, release fence, or both an acquire and release fence.

[INTRO.MULTITHREAD]

1.10 Multi-threaded executions and data races

THREAD SUPPORT

- `std::thread`
- `std::mutex`
- `std::shared_mutex`
- `std::condition_variable`
- `std::future`

ATOMIC OPERATIONS

Atomic: indivisible with respect to all other atomic accesses
to that object

MEMORY ORDER

- relaxed
- consume
- acquire
- release
- acq_rel
- seq_cst

RELAXED

- racing accesses
 - indivisible
- enforce cache coherency
- doesn't guarantee visibility

ATOMIC

- `template<class T> struct atomic;`
- `template<> struct atomic<integral>;`
- `template<class T> struct atomic<T*>;`
- `template<class T> struct atomic<floating-point>; // future`

ATOMIC<T>

- load
- store
- exchange
- compare_exchange_{weak, strong}
- is_lock_free
- static constexpr is_always_lock_free // C++17

ATOMIC<INTEGRAL>

- `fetch_add`
- `fetch_sub`
- `fetch_and`
- `fetch_or`
- `fetch_xor`

ATOMIC<T*>

- fetch_add
- fetch_sub

ATOMIC<FLOATING- POINT>

- `fetch_add`
- `fetch_sub`
- More?

FENCES

- atomic_thread_fence
- atomic_signal_fence

LOCK-FREEDOM

Guaranteed for `std::atomic_flag`

Roughly: is there a compare-exchange for this size?

This is what I'm focusing on.

What does non-lock-free mean?

SIMPLE SINGLE PRODUCER / CONSUMER

```
struct Data { int spam, bacon, eggs; std::atomic<int> ready; };  
void write(Data *d) {  
    d->spam = 0xBEEF;  
    d->bacon = 0xCAFE;  
    d->eggs = 0x0000;  
    d->ready.store(1, std::memory_order_release);  
}  
auto read(Data *d) {  
    while (!d->ready.load(std::memory_order_acquire)) ;  
    return std::make_tuple(d->spam, d->bacon, d->eggs);  
}
```

WE HAVE OUR TOOLS!

Why is the language set up this way?

Provide an abstraction for relevant hardware platforms.

X86-64

load relaxed	MOV
load consume	MOV
load acquire	MOV
load seq_cst	MOV
store relaxed	MOV
store release	MOV
store seq_cst	LOCK XCHG
non-seq_cst fence	# free!
seq_cst fence	MFENCE

POWER

load relaxed	ld
load consume	ld + dependencies
load acquire	ld; cmp; bc; isync
load seq_cst	hwsync; ld; cmp; bc; isync
store relaxed	st
store release	lwsync; st
store seq_cst	hwsync; st
cpxchgr relaxed	_loop: lwarx; cmp; bc _exit; stwcx.; bc _loop; _exit:
non-seq_cst fence	lwsync
seq_cst fence	hwsync

ARMV7

load relaxed	ldr
load consume	ldr + dependencies
load acquire	ldr; dmb
load seq_cst	lrd; dmb
store relaxed	str
store release	dmb; str
store seq_cst	dmb; str; dmb
cmpxchg relaxed	_loop: ldrex; mov 0; teq; strexeq; teq 0; bne _loop
fences	dmb

ARMV8 A64

load relaxed	ldr
load consume	ldr + dependencies
load acquire	ldar
load seq_cst	ldar
store relaxed	str
store release	stlr
store seq_cst	stlr
cmpxchg relaxed	_loop: ldxr; cmp; bne _exit; stxr; cbz _loop; _exit
fences	dmb

ITANIUM

load relaxed	ld.acq
load consume	ld.acq
load acquire	ld.acq
load seq_cst	ld.acq
store relaxed	st.rel
store release	st.rel
store seq_cst	st.rel; mfb
cpxchg acquire	cpxchg.acq
non-seq_cst fence	# free!
seq_cst fence	mf

ALPHA

Oh, one of my favorite (NOT!) pieces of code in the kernel is the implementation of the `smp_read_barrier_depends()` macro, which on every single architecture except for one (alpha) is a no-op.

We have basically 30 or so empty definitions for it, and I think we have something like five uses of it. One of them, I think, is performance critical, and the reason for that macro existing.

What does it do? The semantics is that it's a read barrier between two different reads that we want to happen in order wrt two writes on the writing side (the writing side also has to have a `smp_wmb()` to order those writes). But the reason it isn't a simple read barrier is that the reads are actually causally *dependent*, ie we have code like

```
first_read = read_pointer;
smp_read_barrier_depends();
second_read = *first_read;
```

and it turns out that on pretty much all architectures (except for alpha), the *data*dependency* will already guarantee that the CPU reads the thing in order. And because a read barrier can actually be quite expensive, we don't want to have a read barrier for this case.

But alpha? Its memory consistency is so broken that even the data dependency doesn't actually guarantee cache access order. It's strange, yes. No, it's not that alpha does some magic value prediction and can do the second read without having even done the first read first to get the address. What's actually going on is that the cache itself is unordered, and without the read barrier, you may get a stale version from the cache even if the writes were forced (by the write barrier in the writer) to happen in the right order.

CONSUME FOR ARM

Address dependency rule

```
LDR r1, [r0]  
LDR r2, [r1]
```

```
LDR r1, [r0]  
AND r1, r1, #0  
LDR r2, [r3, r1]
```

Look ma, no barrier!

ARCHITECTURE-SPECIFIC

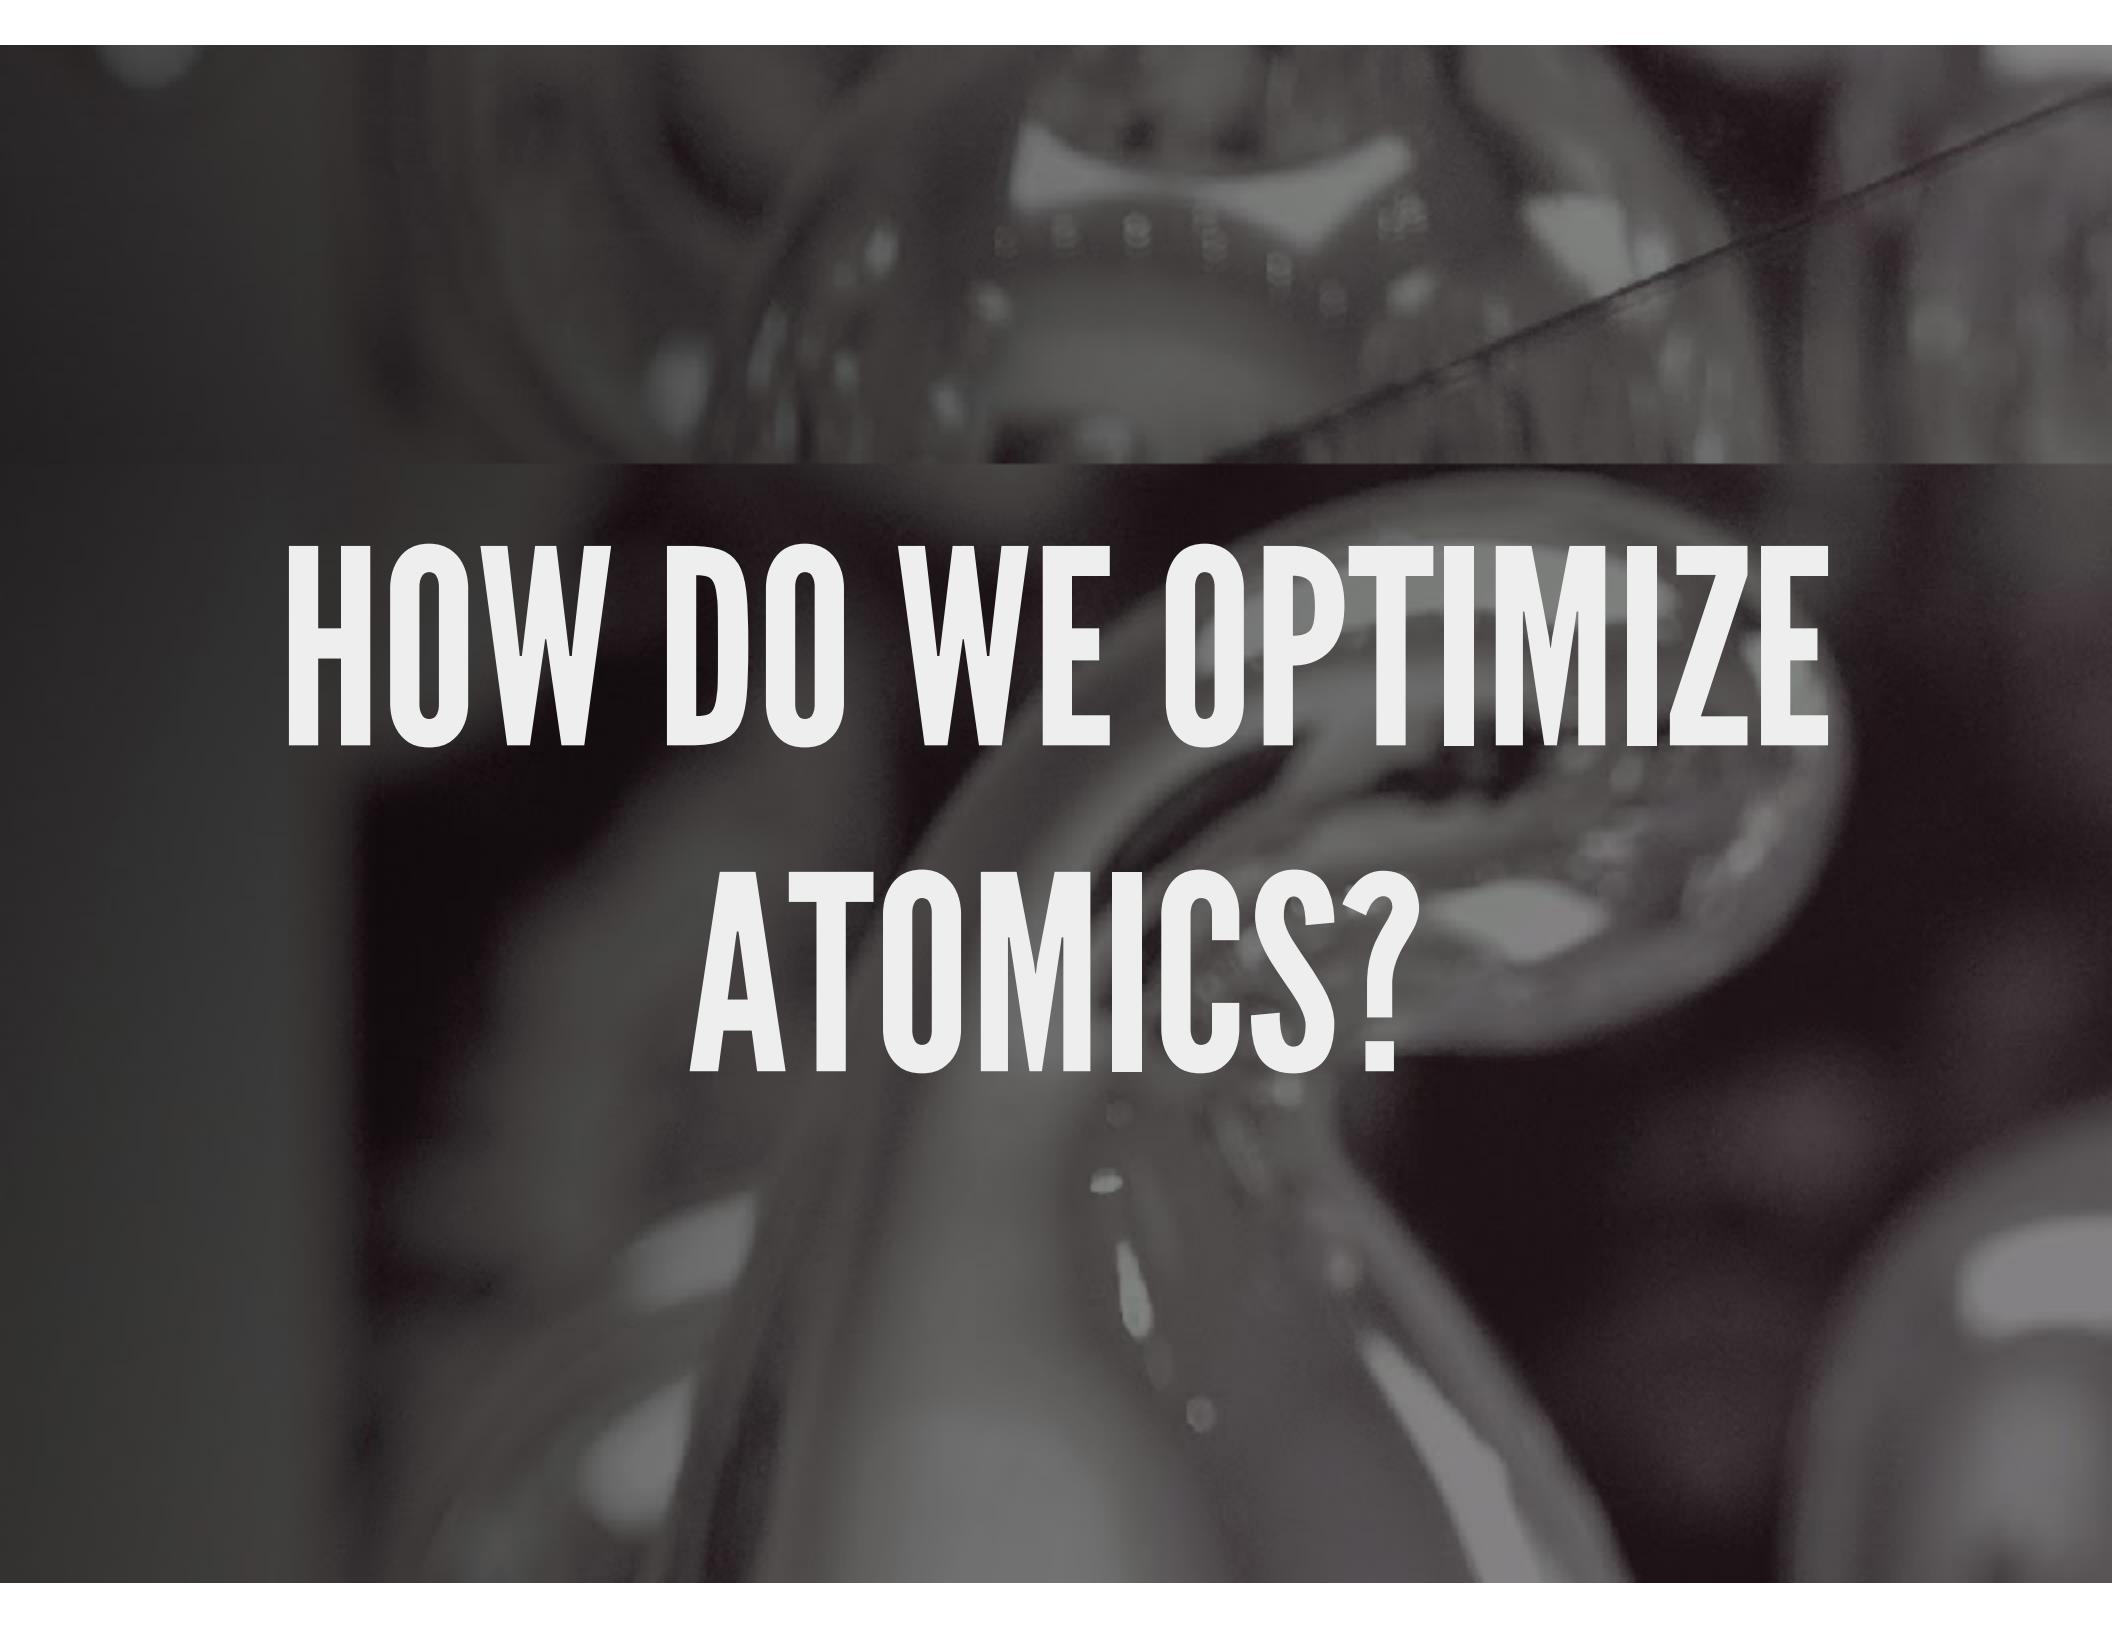
Same architecture, different instructions / performance.

Flags to the rescue: -march, -mcpu, -mattr

- Special instructions
- what's the best spinloop?
- nominally incorrect code
- Workaround CPU bugs

TAKEAWAYS

- C++ exposes hardware portably
- Hardware changes over time
- Assembly operations operation on memory locations
- C++ atomics operate on datastructures



**HOW DO WE OPTIMIZE
ATOMICS?**

AS-IF

more atomic: don't violate forward progress

less atomic: don't add non-benign race which weren't already present

PUT ANOTHER WAY

correct programs must work under all executions an implementation is allowed to create

WHAT CAN THE COMPILER DO?

At least as much as the hardware could do
and maybe more ☺

WHAT DOES HARDWARE DO?

We saw a few examples earlier

SIMPLE EXAMPLE

```
void inc(std::atomic<int> *y) {  
    *y += 1;  
}  
std::atomic<int> x;  
void two() {  
    inc(&x);  
    inc(&x);  
}
```

```
std::atomic<int> x;  
void two() {  
    x += 2;  
}
```

SIMILAR EXAMPLE

```
std::atomic<int> x;
void inc(int val) {
    x += 1;
    x += val;
}
```

```
_Z3inci:
    lock incl x(%rip)
    lock addl %edi, x(%rip)
```

OPPORTUNITIES THROUGH INLINING

```
template<typename T>
bool silly(std::atomic<T> *x, T expected, T desired) {
    x->compare_exchange_strong(expected, desired); // Inlined.
    return expected == desired;
}
```

```
template<typename T>
bool silly(std::atomic<T> *x, T expected, T desired) {
    return x->compare_exchange_strong(expected, desired);
}
```

A TRICKY ONE

Works for any memory order but release and acq_rel

```
template<typename T>
bool optme(std::atomic<T> *x, T desired) {
    T expected = desired;
    return x->compare_exchange_strong(expected, desired
        std::memory_order_seq_cst, std::memory_order_relaxed);
}
```

```
template<typename T>
bool optme(std::atomic<T> *x, T desired) {
    return x->load(std::memory_order_seq_cst) == desired; // †
}
```

† compiler: mark transformed load as *release sequence* †

‡ as defined in section 1.10 of the C++ standard

STRONGER, FASTER

```
template<typename T>
T optmetoo(std::atomic<T> *x, T y) {
    T z = x->load();
    x->store(y);
    return z;
}
```

```
template<typename T>
T optmetoo(std::atomic<T> *x, T y) {
    return x->exchange(y);
}
```

BETTER?

INTERESTING PROBLEM...



READ ; MODIFY ; WRITE

Hogwild!: A Lock-Free Approach to Parallelizing Stochastic
Gradient Descent

RACY

```
float y;  
void g(float a) { y += a; }
```

```
addss    y(%rip), %xmm0  
movss    %xmm0, y(%rip)
```

VOLATILE

```
volatile float yv;  
void gv(float a) { yv += a; }
```

```
addss    yv(%rip), %xmm0  
movss    %xmm0, yv(%rip)
```

CORRECT?

```
std::atomic<float> x;
void f(float a) {
    x.store(x.load(std::memory_order_relaxed) + a,
            std::memory_order_relaxed);
}
```

```
movl    x(%rip), %eax
movd    %eax, %xmm1
addss  %xmm0, %xmm1
movd    %xmm1, %eax
movl    %eax, x(%rip)
```

```
addss  x(%rip), %xmm0
movss  %xmm0, x(%rip)
```

Optimization FTW!

MOAR!

Inlining and constant propagation

```
atomic<T>::fetch_and(~(T)0) →  
atomic<T>::load()
```

Same for `fetch_or(0)` and `fetch_xor(0)`

```
atomic<T>::fetch_and(0) →  
atomic<T>::store(0)
```

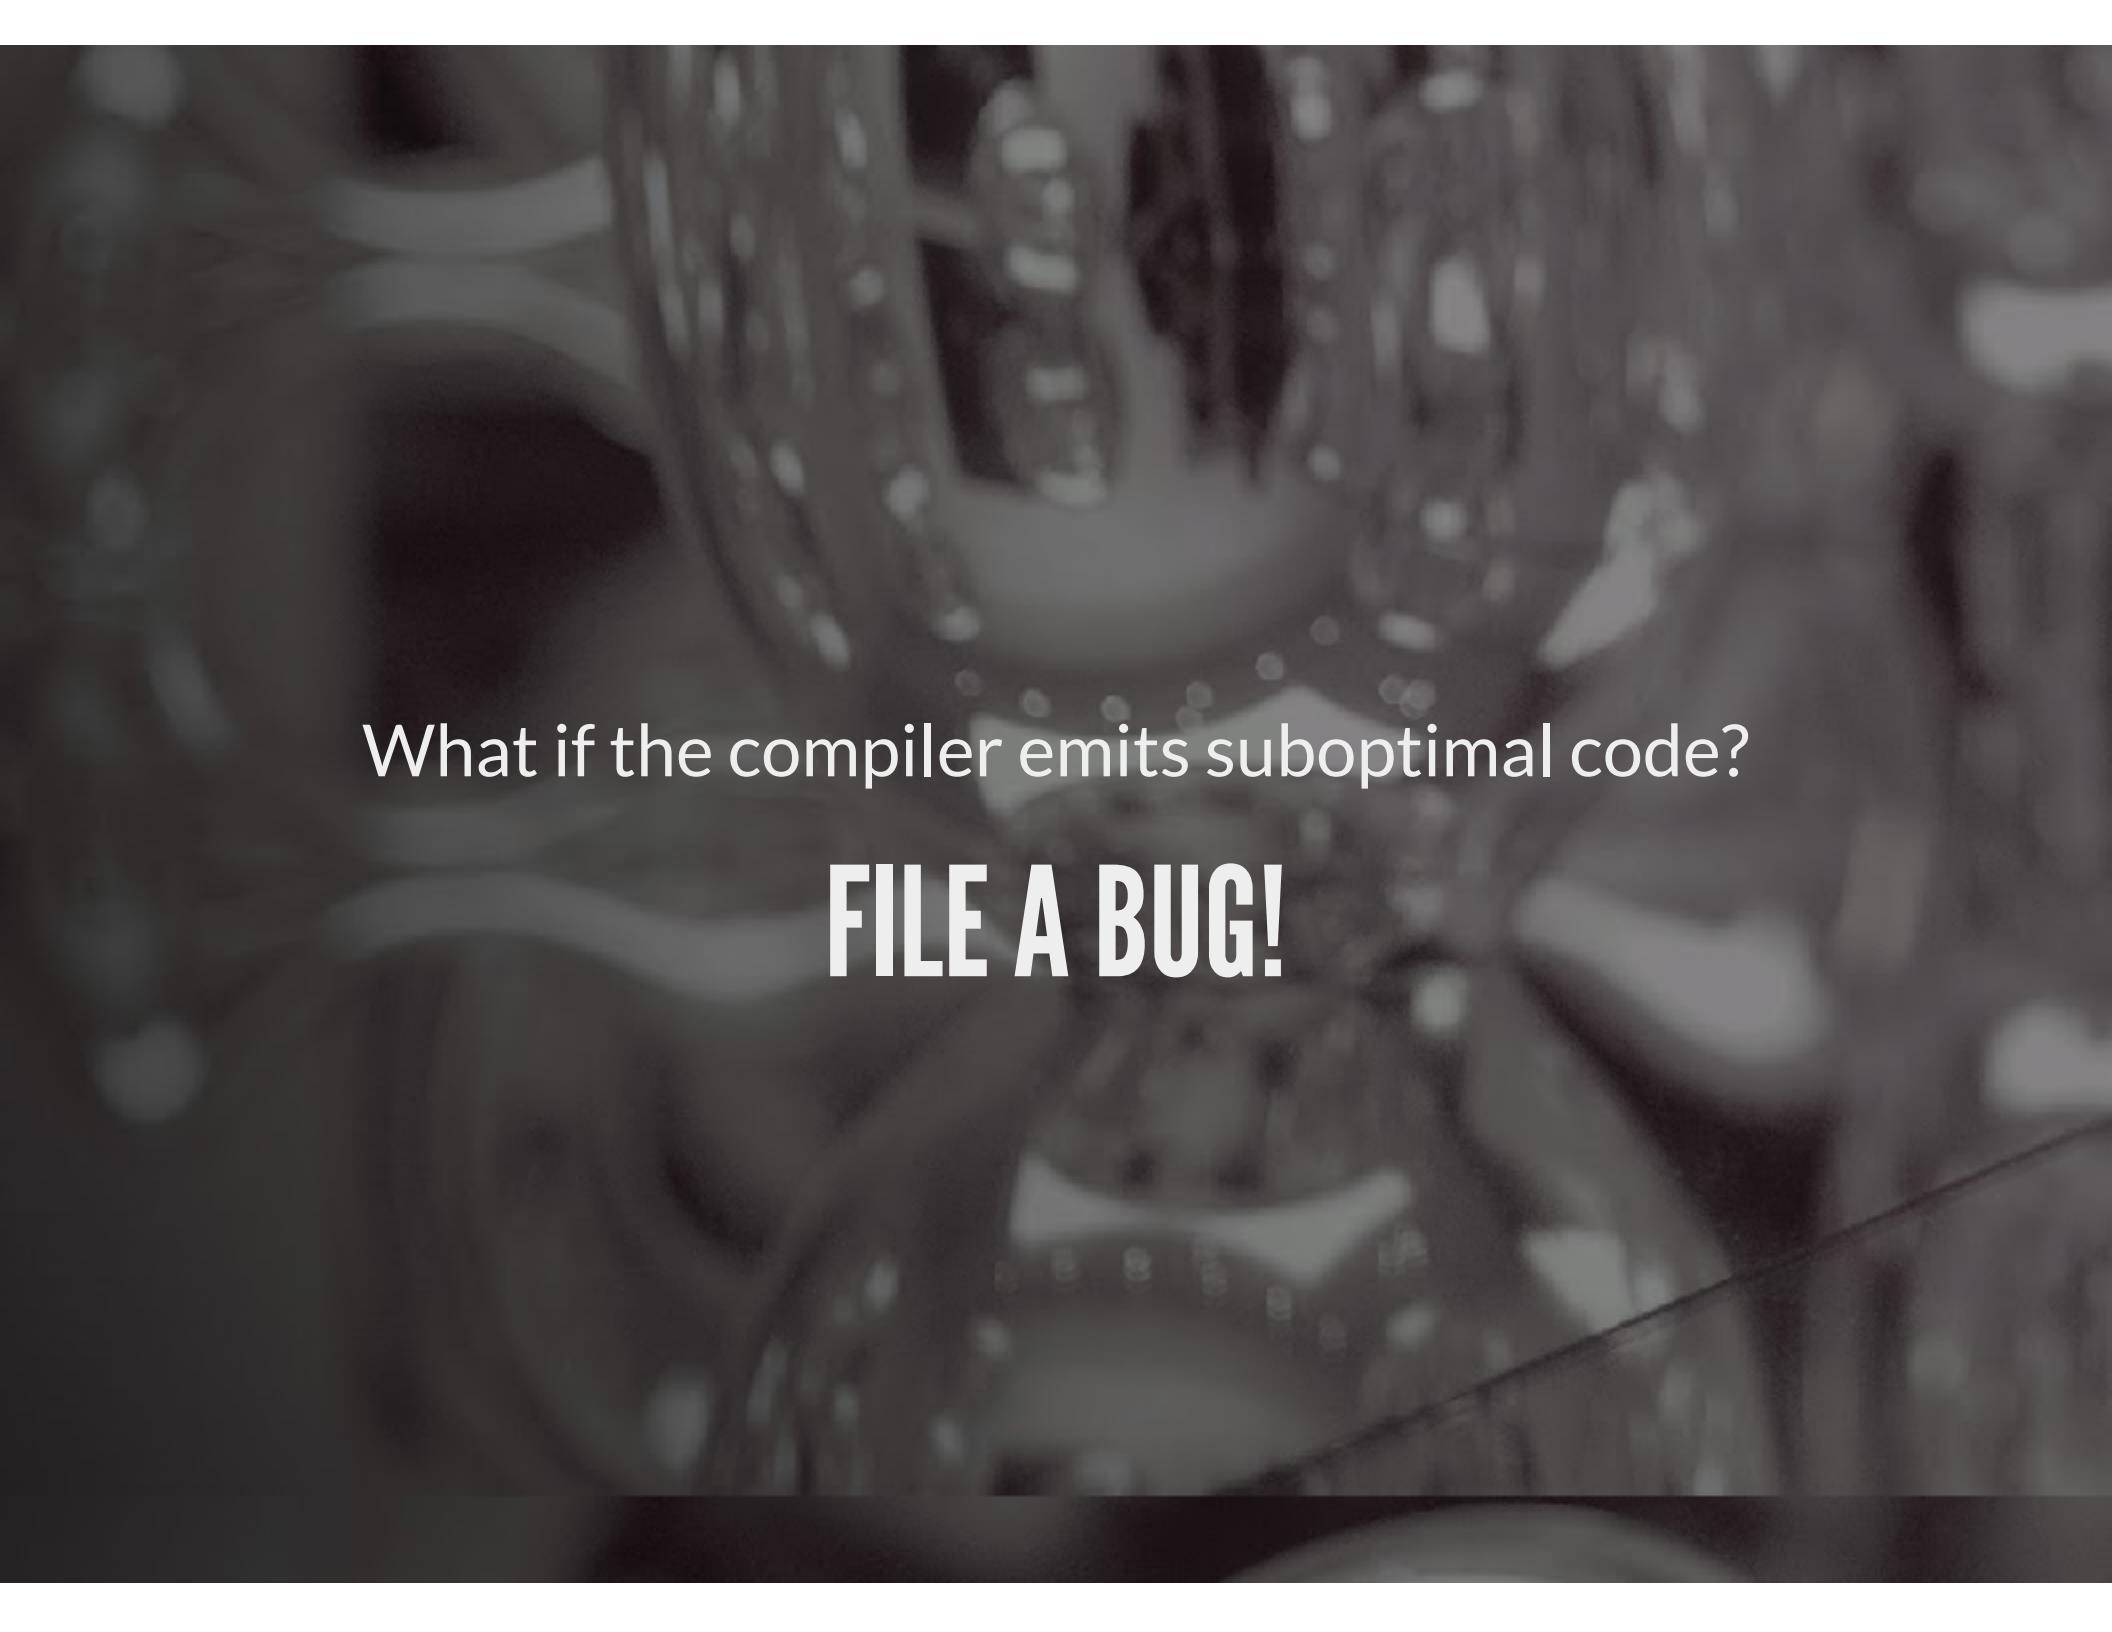
TAKEAWAY

If simple things are hard...

...is your inline assembly correct?

...will it remain correct?

Do you trust your compiler?



What if the compiler emits suboptimal code?

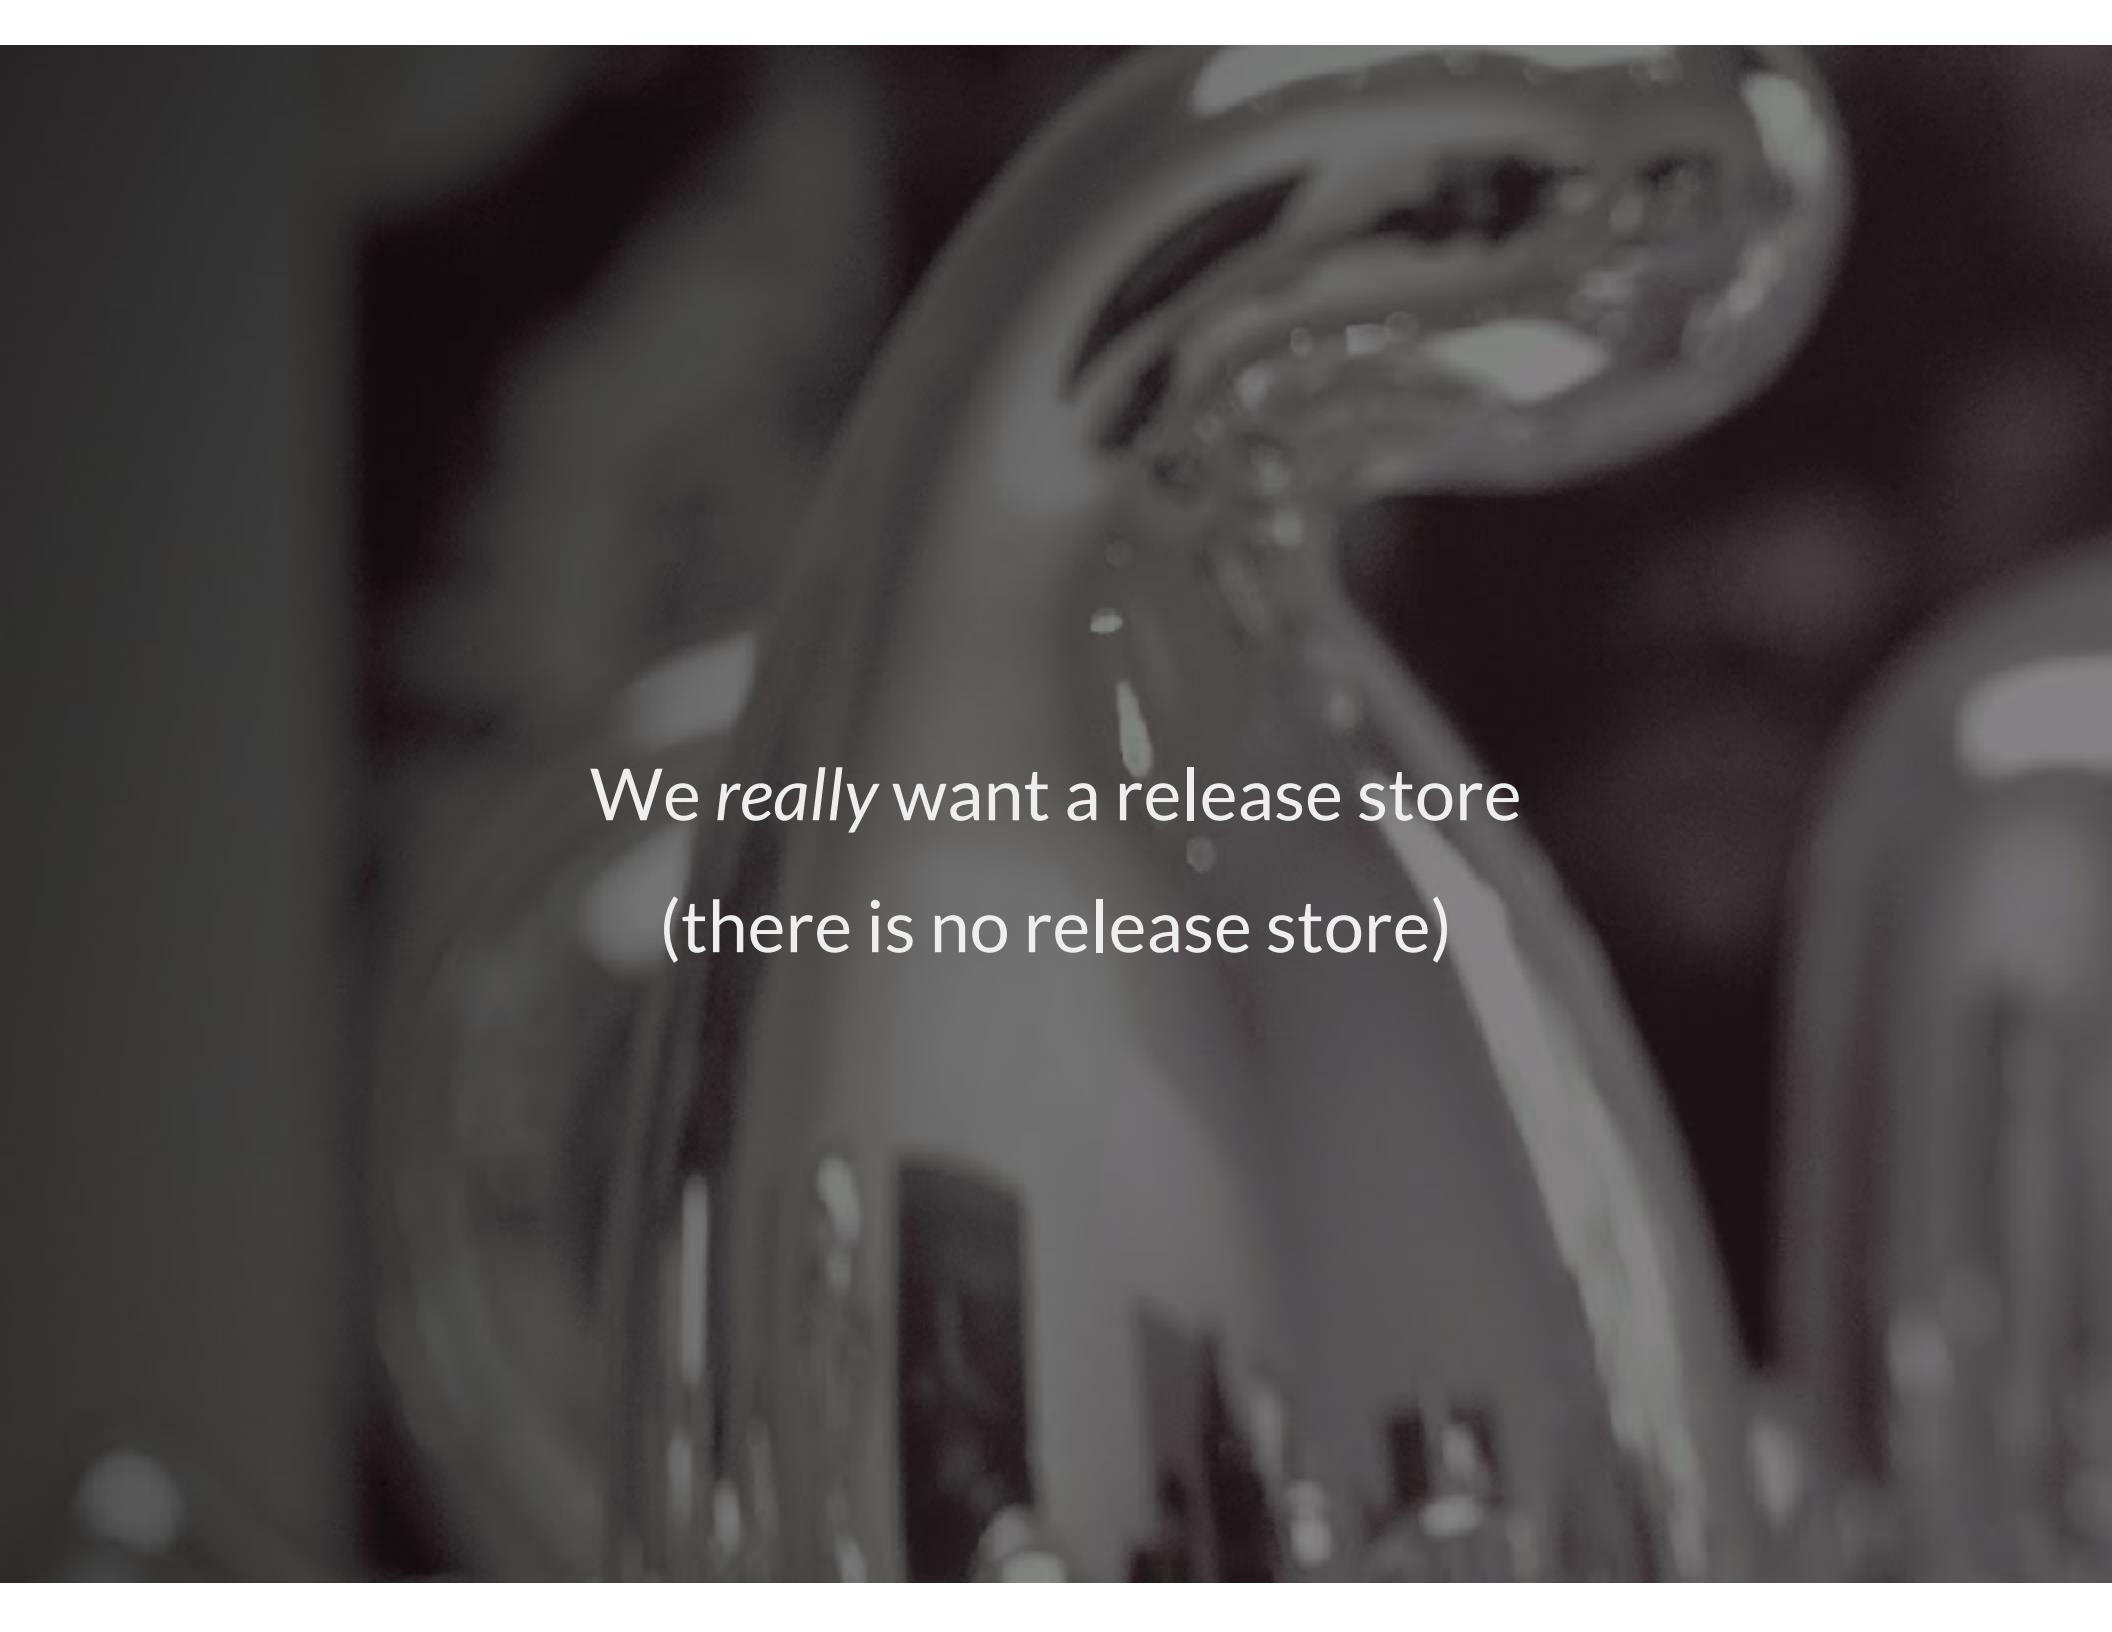
FILE A BUG!

SEQUENCE LOCK

- Get ticket number
- Get the data
- Check the ticket again
 - If odd: write was happening
 - If different: write occurred
 - If same: no write occurred, data is good
- If data isn't good: try again

```
std::tuple<T, T> reader() {
    T d1, d2; unsigned seq0, seq1;
    do {
        seq0 = seq.load(std::memory_order_acquire);
        d1 = data1.load(std::memory_order_relaxed);
        d2 = data2.load(std::memory_order_relaxed);
        std::atomic_thread_fence(std::memory_order_acquire);
        seq1 = seq.load(std::memory_order_relaxed);
    } while (seq0 != seq1 || seq0 & 1);
    return {d1, d2};
}

void writer(T d1, T d2) {
    unsigned seq0 = seq.load(std::memory_order_relaxed);
    seq.store(seq0 + 1, std::memory_order_relaxed);
    data1.store(d1, std::memory_order_release);
    data2.store(d2, std::memory_order_release);
    seq.store(seq0 + 2, std::memory_order_release);
}
```



We *really* want a release store
(there is no release store)

READ-DON'T-MODIFY-WRITE

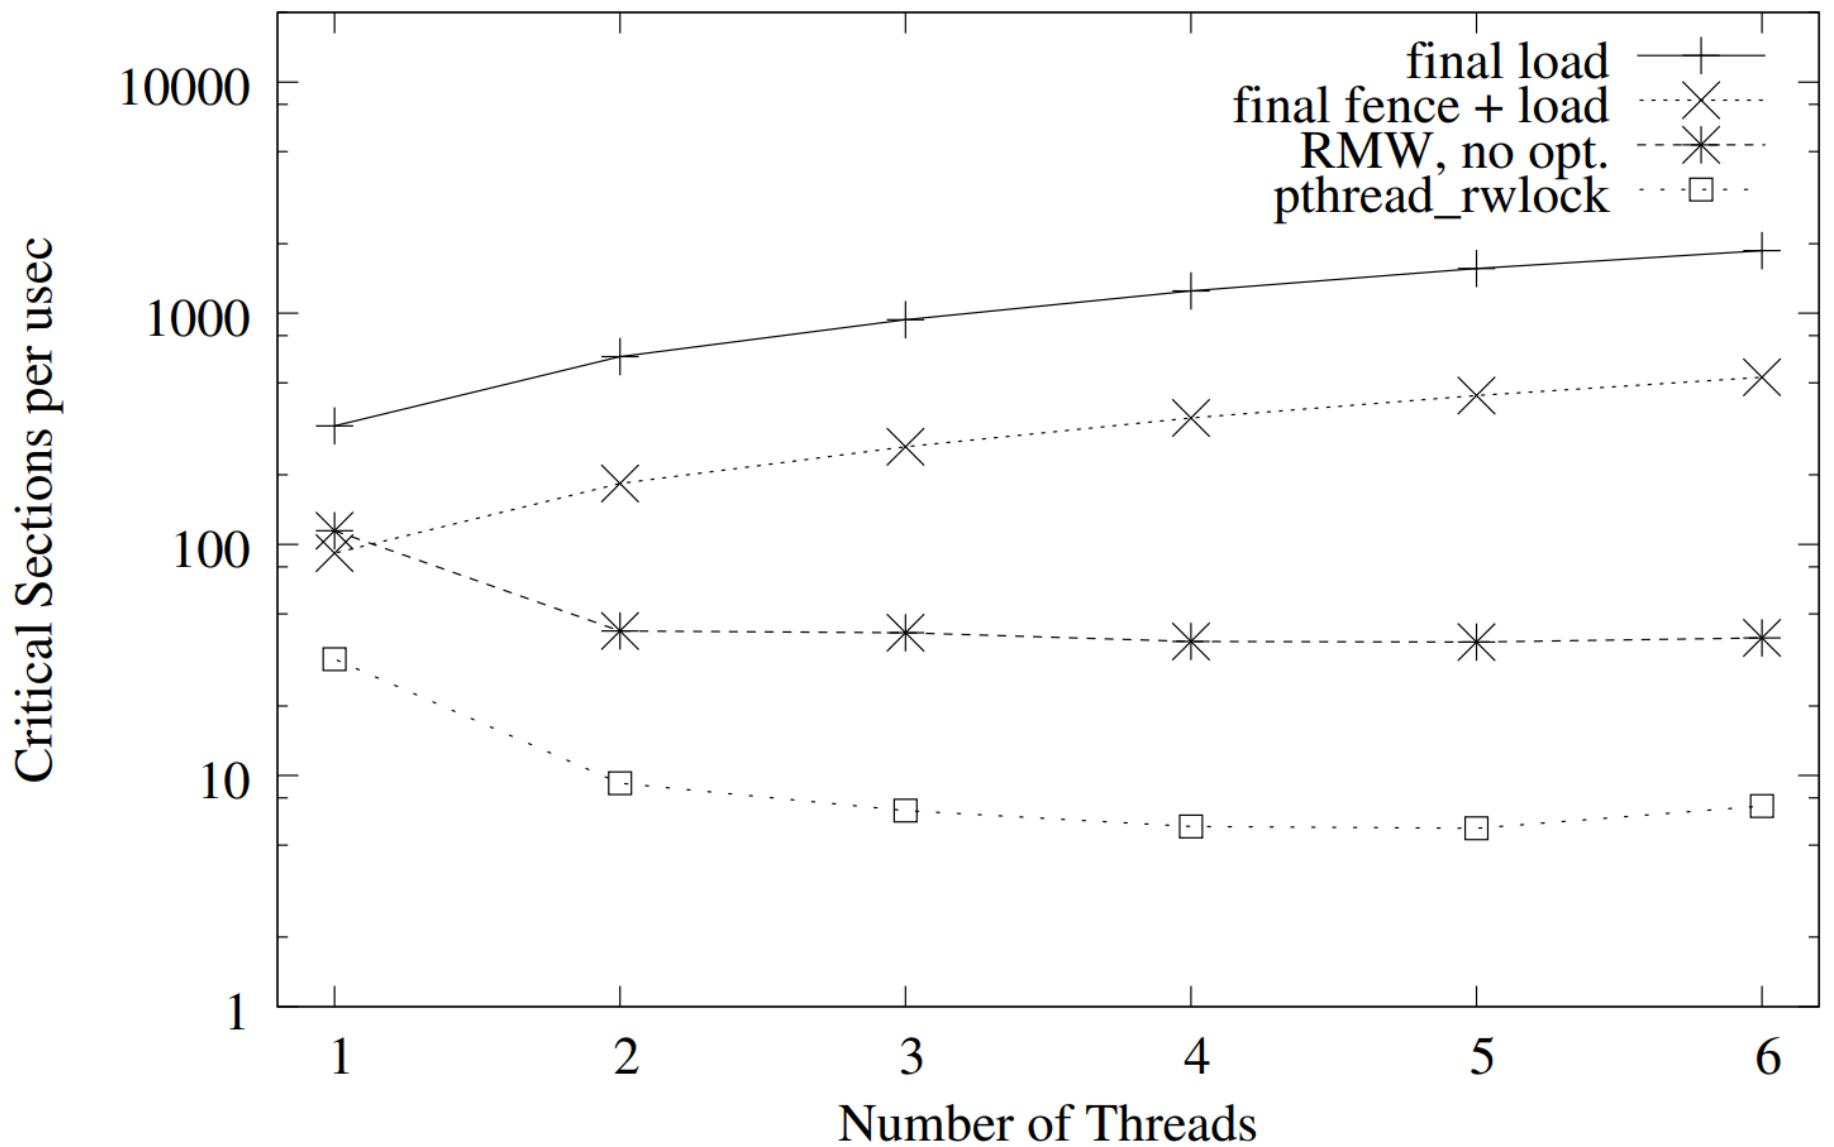
```
T d1, d2;
unsigned seq0, seq1;
do {
    seq0 = seq.load(std::memory_order_acquire);
    d1 = data1.load(std::memory_order_relaxed);
    d2 = data2.load(std::memory_order_relaxed);
    seq1 = seq.fetch_add(0, std::memory_order_release);
} while (seq0 != seq1 || seq0 & 1);
```

```
mov    0x200a76(%rip),%edx      # seq
mov    0x200a74(%rip),%eax      # data1
mov    0x200a72(%rip),%ecx      # data2
# acquire fence
mov    0x200a64(%rip),%esi      # seq
```

```
mov    0x2004c6(%rip),%ecx      # seq
mov    0x2004bc(%rip),%esi      # data1
xor    %edx,%edx
mov    0x2004af(%rip),%r8d      # data2
lock xadd %edx,0x2004af(%rip)    # seq RdMW
```

```
mov    0x200a46(%rip),%edx      # seq
mov    0x200a44(%rip),%eax      # data1
mov    0x200a42(%rip),%ecx      # data2
mfence                         # optimized RdMW!
mov    0x200a31(%rip),%esi      # seq
```

WHAT'S BETTER?



"NORMAL" COMPILER OPTIMIZATIONS

Dead store elimination?

Strength reduction?

RELAXED OPTIMIZATION

```
std::atomic<int> x, y;
void relaxed() {
    x.fetch_add(1, std::memory_order_relaxed);
    y.fetch_add(1, std::memory_order_relaxed);
    x.fetch_add(1, std::memory_order_relaxed);
    y.fetch_add(1, std::memory_order_relaxed);
}
```

```
std::atomic<int> x, y;
void relaxed() {
    x.fetch_add(2, std::memory_order_relaxed);
    y.fetch_add(2, std::memory_order_relaxed);
}
```

PROGRESS BAR

```
atomic<int> progress(0);
f() {
    for (i = 0; i < 1000000; ++i) {
        // Heavy work, no shared memory...
        ++progress;
    }
}
```

```
atomic<int> progress(0);
f() {
    int my_progress = 0; // register allocated
    for (i = 0; i < 1000000; ++i) {
        // Heavy work, no shared memory...
        ++my_progress;
    }
    progress += my_progress;
}
```

DISABLING REORDERING / FUSING?

- `asm volatile(":::";"memory");` // eek!
- `std::atomic_signal_fence();` // wat?
- Use `volatile`? **NO!**
- Don't use `relaxed`?
- Synchronize-with?

MOAR!

- Tag "non-atomic" functions, optimize around them
- Interference-free regions
- Optimize fence positioning
- Non-escaping thread local (TLS, stack)
- etc.

STD::MUTEX FOR SANITY

- easier to use correctly
- in theory...
 - it could still get optimized
 - lock: acquire; unlock: release
- pthread and kernel call makes this hard
- std::synchronic to the rescue!



STD::MUTEX

Knows better

VOLATILE

Initial motivation: `getChar()` on PDP-11, reading memory-mapped KBD_STAT I/O register

K&R: The purpose of `volatile` is to force an implementation to suppress optimization that could otherwise occur

WHAT DOES THAT MEAN?

- Prevent load / store motion?
- Prevent arithmetic motion?
- Hardware fence? x86 MFENCE? ARM DMB SY?
- Shared memory?

A black and white close-up photograph of a person's face. The person has dark hair and is looking off-camera with a contemplative expression. Their right hand is resting against their cheek, with fingers partially hidden in their hair. The lighting is soft, creating a moody and intimate atmosphere.

IT GETS BETTER

XTENSA GCC OPTIONS

- mserialize-volatile
- mnoserialize-volatile

When this option is enabled, GCC inserts MEMW instructions before volatile memory references to guarantee sequential consistency.

MSVC

You can use the `/volatile` compiler switch to modify how the compiler interprets this keyword. If a struct member is marked as `volatile`, then `volatile` is propagated to the whole structure. If a structure does not have a length that can be copied on the current architecture by using one instruction, `volatile` may be completely lost on that structure.

clang emulates this with `-fms-volatile`

WHAT DOES IT MEAN?

-＼(ツ)／-

- can tear
- each byte touched exactly once
- no machine-level ordering
- enforces abstract machine ordering with other volatile
- can't duplicate
- can't elide

CAN VOLATILE BE LOCK-FREE?

↖_(ツ)_↗

CAN VOLATILE DO RMW?

No

FUN VOLATILE FACT

Is this undefined behavior?

```
volatile int foo;  
void bar() { int baz; while (1) baz = foo; }
```

The implementation may assume that any thread will eventually do one of the following:

- terminate,
- make a call to a library I/O function,
- access or modify a volatile object, or
- perform a synchronization operation or an atomic operation.

ARE THERE PORTABLE VOLATILE USES?

- mark local variables in the same scope as a `setjmp`,
preserve value across `longjmp`
- **externally modified**: physical memory mapped in multiple
virtual addresses
- `volatile sigatomic_t` may be used to communicate
with a signal handler in the same thread

VOLATILE ATOMIC

what's that?

NON-TEMPORALS

WHAT DO COMPILERS DO?

by compilers I mean clang

STANDARD LIBRARY

C's `_Atomic` ↔ C++'s `std::atomic`

`_sync_*` builtins

`_atomic_*` builtins

FRONT-END

does little besides inlining constant `memory_order_*`

MIDDLE-END

- load atomic [volatile] <ty>, <ty>*<pointer> [singlethread] <ordering>, align <alignment>
- store atomic [volatile] <ty> <value>, <ty>*<pointer> [singlethread] <ordering>, align <alignment>

MIDDLE-END

- atomicrmw [volatile] <operation> <ty>* <pointer>, <ty> <value> [singlethread] <ordering>
- cmpxchg [weak] [volatile] <ty>* <pointer>, <ty> <cmp>, <ty> <new> [singlethread] <success ordering> <failure ordering>
- fence [singlethread] <ordering>

IR

isSimple, isVolatile, isAtomic: steer clear

MACHINE IR

similar to actual ISA instructions



NOW WHAT?

STANDARDS COMMITTEE

Assume optimizations occur, encourage them.

Standardize common practice, enable to-the-metal optimizations.

More libraries: easy to use concurrency & parallelism; hard to get it wrong.

DEVELOPERS

Drop assembly

File bugs

Talk to standards committee.

Use tooling: ThreadSanitizer.

HARDWARE VENDORS

Showcase your hardware's strengths.

COMPILER WRITERS

Get back to work!

SANE COMPILERS SHOULD OPTIMIZE ATOMICS

and you should use atomics

github.com/jfbastien/no-sane-compiler

@jfbastien